

# Lasso Regression

*dhruv patel BDA 17162121014*

## Loading Libraries

```
In [2]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4
        5 from sklearn.datasets import load_breast_cancer
        6 from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
        7
        8 from sklearn.linear_model import LinearRegression
        9 from sklearn.linear_model import Lasso, Ridge
       10
       11 from sklearn.metrics import accuracy_score
```

```
In [3]: 1 data = load_breast_cancer() # Loading data
        2 data
```

```
Out[3]: {'data': array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01, 4.601e-01,
    1.189e-01],
    [2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
    8.902e-02],
    [1.969e+01, 2.125e+01, 1.300e+02, ..., 2.430e-01, 3.613e-01,
    8.758e-02],
    ...,
    [1.660e+01, 2.808e+01, 1.083e+02, ..., 1.418e-01, 2.218e-01,
    7.820e-02],
    [2.060e+01, 2.933e+01, 1.401e+02, ..., 2.650e-01, 4.087e-01,
    1.240e-01],
    [7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 2.871e-01,
    7.039e-02]]),
  'target': array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
    1, 1,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
    0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0,
    1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0,
    1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1,
    1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0,
    0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
    1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1,
    0, 0,
    0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0,
    1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1,
    1,
    1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
    0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1,
    1,
    1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1,
    1,
    1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0,
    0,
    0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0,
    0,
    0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0,
    0,
    1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
    1,
    1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0,
    0,
    1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1,
    1,
    1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
    1,
    1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
    1,
    1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1]),
  'frame': None,
  'target_names': array(['malignant', 'benign'], dtype='<U9'),
  'DESCR': '.. _breast_cancer_dataset:\n\nBreast cancer wisconsin (diagnostic) d
ataset\n-----\n\n**Data Set Characterist
ics:**\n\n    :Number of Instances: 569\n\n    :Number of Attributes: 30 numeri
c, predictive attributes and the class\n\n    :Attribute Information:\n
- radius (mean of distances from center to points on the perimeter)\n
- texture (standard deviation of gray-scale values)\n
- perimeter\n
- area\n
- smoothness (local variation in radius lengths)\n
- compactness (perimeter^2 / area - 1.0)\n
- concavity (severity of concave p
ortions of the contour)\n
- concave points (number of concave portions o
f the contour)\n
- symmetry\n
- fractal dimension ("coastline app
roximation" - 1)\n\n    The mean, standard error, and "worst" or largest (m
ean of the three\n
    worst/largest values) of these features were computed
```

```

for each image,\n          resulting in 30 features. For instance, field 0 is Me
an Radius, field\n          10 is Radius SE, field 20 is Worst Radius.\n\n
- class:\n          - WDBC-Malignant\n          - WDBC-Benign\n\n
:Summary Statistics:\n\n          =====
\n          Min    Max\n          =====
===== \n          radius (mean):          6.
981 28.11\n          texture (mean):          9.71 39.28\n          perimet
er (mean):          43.79 188.5\n          area (mean):
143.5 2501.0\n          smoothness (mean):          0.053 0.163\n          comp
actness (mean):          0.019 0.345\n          concavity (mean):
0.0 0.427\n          concave points (mean):          0.0 0.201\n          symme
try (mean):          0.106 0.304\n          fractal dimension (mean):
0.05 0.097\n          radius (standard error):          0.112 2.873\n          textu
re (standard error):          0.36 4.885\n          perimeter (standard error):
0.757 21.98\n          area (standard error):          6.802 542.2\n          smoot
hness (standard error):          0.002 0.031\n          compactness (standard erro
r):          0.002 0.135\n          concavity (standard error):          0.0 0.39
6\n          concave points (standard error):          0.0 0.053\n          symmetry (standa
rd error):          0.008 0.079\n          fractal dimension (standard error):
0.001 0.03\n          radius (worst):          7.93 36.04\n          textur
e (worst):          12.02 49.54\n          perimeter (worst):
50.41 251.2\n          area (worst):          185.2 4254.0\n          smoo
thness (worst):          0.071 0.223\n          compactness (worst):
0.027 1.058\n          concavity (worst):          0.0 1.252\n          conca
ve points (worst):          0.0 0.291\n          symmetry (worst):
0.156 0.664\n          fractal dimension (worst):          0.055 0.208\n          =====
===== \n\n          :Missing Attribute Value
s: None\n\n          :Class Distribution: 212 - Malignant, 357 - Benign\n\n          :Creat
or: Dr. William H. Wolberg, W. Nick Street, Olvi L. Mangasarian\n\n          :Donor:
Nick Street\n\n          :Date: November, 1995\n\nThis is a copy of UCI ML Breast Can
cer Wisconsin (Diagnostic) datasets.\nhttps://goo.gl/U2Uwz2\n\nFeatures are com
puted from a digitized image of a fine needle\naspirate (FNA) of a breast mass.
They describe\ncharacteristics of the cell nuclei present in the image.\n\nSepa
rating plane described above was obtained using\nMultisurface Method-Tree (MSM-
T) [K. P. Bennett, "Decision Tree\nConstruction Via Linear Programming." Procee
dings of the 4th\nMidwest Artificial Intelligence and Cognitive Science Societ
y,\npp. 97-101, 1992], a classification method which uses linear\nprogramming t
o construct a decision tree. Relevant features\nwere selected using an exhaust
ive search in the space of 1-4\nfeatures and 1-3 separating planes.\n\nThe actu
al linear program used to obtain the separating plane\nin the 3-dimensional spa
ce is that described in:\n[K. P. Bennett and O. L. Mangasarian: "Robust Linear
\nProgramming Discrimination of Two Linearly Inseparable Sets",\nOptimization M
ethods and Software 1, 1992, 23-34].\n\nThis database is also available through
the UW CS ftp server:\n\nftp ftp.cs.wisc.edu\ncd math-prog/cpo-dataset/machine-
learn/WDBC/\n\n.. topic:: References\n\n          - W.N. Street, W.H. Wolberg and O.L.
Mangasarian. Nuclear feature extraction \n          for breast tumor diagnosis. IS&
T/SPIE 1993 International Symposium on \n          Electronic Imaging: Science and T
echnology, volume 1905, pages 861-870,\n          San Jose, CA, 1993.\n          - O.L. Man
gasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and \n          prog
nosis via linear programming. Operations Research, 43(4), pages 570-577, \n
July-August 1995.\n          - W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machin
e learning techniques\n          to diagnose breast cancer from fine-needle aspirate
s. Cancer Letters 77 (1994) \n          163-171.',
'feature_names': array(['mean radius', 'mean texture', 'mean perimeter', 'mean
area',
                        'mean smoothness', 'mean compactness', 'mean concavity',
                        'mean concave points', 'mean symmetry', 'mean fractal dimension',

```

```
'radius error', 'texture error', 'perimeter error', 'area error',
'smoothness error', 'compactness error', 'concavity error',
'concave points error', 'symmetry error',
'fractal dimension error', 'worst radius', 'worst texture',
'worst perimeter', 'worst area', 'worst smoothness',
'worst compactness', 'worst concavity', 'worst concave points',
'worst symmetry', 'worst fractal dimension'], dtype='<U23'),
'filename': 'C:\\Users\\dhruv\\anaconda3\\envs\\ml\\lib\\site-packages\\sklear
n\\datasets\\data\\breast_cancer.csv'}
```

```
In [4]: 1 data.target_names
```

```
Out[4]: array(['malignant', 'benign'], dtype='<U9')
```

```
In [5]: 1 data.feature_names
```

```
Out[5]: array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
'mean smoothness', 'mean compactness', 'mean concavity',
'mean concave points', 'mean symmetry', 'mean fractal dimension',
'radius error', 'texture error', 'perimeter error', 'area error',
'smoothness error', 'compactness error', 'concavity error',
'concave points error', 'symmetry error',
'fractal dimension error', 'worst radius', 'worst texture',
'worst perimeter', 'worst area', 'worst smoothness',
'worst compactness', 'worst concavity', 'worst concave points',
'worst symmetry', 'worst fractal dimension'], dtype='<U23')
```

```
In [6]: 1 df = pd.DataFrame(data.data, columns=data.feature_names)
```

```
In [7]: 1 df.head()
```

```
Out[7]:
```

	mean radius	mean texture	mean concave points	mean symmetry	mean fractal dimension	...	worst radius	worst texture	worst perimeter	worst area	worst smoothness
27760	0.3001	0.14710	0.2419	0.07871	...	25.38	17.33	184.60	2019.0	0.1622	
07864	0.0869	0.07017	0.1812	0.05667	...	24.99	23.41	158.80	1956.0	0.1238	
15990	0.1974	0.12790	0.2069	0.05999	...	23.57	25.53	152.50	1709.0	0.1444	
28390	0.2414	0.10520	0.2597	0.09744	...	14.91	26.50	98.87	567.7	0.2098	
13280	0.1980	0.10430	0.1809	0.05883	...	22.54	16.67	152.20	1575.0	0.1374	


In [8]:

1 df.describe()

Out[8]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540
75%	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800

8 rows × 30 columns



In [9]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                          569 non-null    float64
1   mean texture                         569 non-null    float64
2   mean perimeter                      569 non-null    float64
3   mean area                          569 non-null    float64
4   mean smoothness                    569 non-null    float64
5   mean compactness                   569 non-null    float64
6   mean concavity                     569 non-null    float64
7   mean concave points                569 non-null    float64
8   mean symmetry                      569 non-null    float64
9   mean fractal dimension              569 non-null    float64
10  radius error                       569 non-null    float64
11  texture error                      569 non-null    float64
12  perimeter error                    569 non-null    float64
13  area error                        569 non-null    float64
14  smoothness error                   569 non-null    float64
15  compactness error                  569 non-null    float64
16  concavity error                    569 non-null    float64
17  concave points error               569 non-null    float64
18  symmetry error                     569 non-null    float64
19  fractal dimension error            569 non-null    float64
20  worst radius                       569 non-null    float64
21  worst texture                      569 non-null    float64
22  worst perimeter                    569 non-null    float64
23  worst area                        569 non-null    float64
24  worst smoothness                   569 non-null    float64
25  worst compactness                  569 non-null    float64
26  worst concavity                    569 non-null    float64
27  worst concave points               569 non-null    float64
28  worst symmetry                     569 non-null    float64
29  worst fractal dimension             569 non-null    float64
dtypes: float64(30)
memory usage: 133.5 KB
```

In [10]: 1 X = data.data

In [11]: 1 y = data.target

In [12]: 1 X\_train,X\_test,y\_train,y\_test = train\_test\_split(X,y,test\_size=0.3,random\_st

## Building Models

In [13]: 1 lr = LinearRegression()

```
In [14]: 1 lr.fit(X_train,y_train)
```

```
Out[14]: LinearRegression()
```

```
In [15]: 1 lr.score(X_test,y_test)  ## test score for linear regression
```

```
Out[15]: 0.739386998952047
```

```
In [16]: 1 lasso_1e1 = Lasso(alpha=0.001 , max_iter = 10e6)
         2
```

```
In [17]: 1 lasso_1e1.fit(X_train,y_train)
```

```
Out[17]: Lasso(alpha=0.001, max_iter=10000000.0)
```

```
In [19]: 1 lasso_1e1.score(X_test,y_test)  ## we can see that our score is quite low be
```

```
Out[19]: 0.7033245684677331
```

```
In [20]: 1 ### trying to search best ALpha using GridSearchcv
```

```
In [21]: 1 lasso = Lasso()
2 param_vals = [1e-20, 1e-10, 1e-5, 1e-1, 1, 10, 100, 20]
3 params = {'alpha': param_vals, "max_iter": [10e5]}
4
5 lassso_cv = GridSearchCV(lasso, params, scoring="neg_mean_squared_error", cv=5,
6
7 lassso_cv.fit(X_train, y_train)
```

C:\Users\dhruv\anaconda3\envs\ml\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:529: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 7.5364956176210125, tolerance: 0.007471698113207546

model = cd\_fast.enet\_coordinate\_descent(  
C:\Users\dhruv\anaconda3\envs\ml\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:529: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 8.030038076620464, tolerance: 0.00742138364779874

model = cd\_fast.enet\_coordinate\_descent(  
C:\Users\dhruv\anaconda3\envs\ml\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:529: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 8.172295272619992, tolerance: 0.007446855345911951

model = cd\_fast.enet\_coordinate\_descent(  
C:\Users\dhruv\anaconda3\envs\ml\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:529: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 8.426045348810645, tolerance: 0.007485893416927903

model = cd\_fast.enet\_coordinate\_descent(  
C:\Users\dhruv\anaconda3\envs\ml\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:529: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 7.92288383145575, tolerance: 0.007460815047021947

model = cd\_fast.enet\_coordinate\_descent(

```
Out[21]: GridSearchCV(cv=5, estimator=Lasso(),
                    param_grid={'alpha': [1e-20, 1e-10, 1e-05, 0.1, 1, 10, 100, 20],
                                'max_iter': [1000000.0]},
                    scoring='neg_mean_squared_error')
```

```
In [22]: 1 lassso_cv.best_params_
```

```
Out[22]: {'alpha': 1e-05, 'max_iter': 1000000.0}
```

```
In [23]: 1 lassso_cv.best_score_
```

```
Out[23]: -0.062488634237653815
```

```
In [24]: 1 lassso_cv.best_estimator_
```

```
Out[24]: Lasso(alpha=1e-05, max_iter=1000000.0)
```



In [25]: 1 lasso\_cv.cv\_results\_

```
Out[25]: {'mean_fit_time': array([3.58654497e+01, 3.28113751e+00, 1.23047647e+00, 4.6011
9247e-03,
        1.00030899e-03, 8.00228119e-04, 4.00066376e-04, 8.00132751e-04]),
'std_fit_time': array([5.39236675e-01, 4.65809269e-01, 2.15030476e-01, 3.32364
010e-03,
        3.87384339e-07, 4.00114329e-04, 4.89979335e-04, 4.00066575e-04]),
'mean_score_time': array([0.00020013, 0.00020003, 0.00060024, 0.00040002, 0.00
019999,
        0.00020013, 0.          , 0.          ]),
'std_score_time': array([0.00040026, 0.00040007, 0.0004901 , 0.00048992, 0.000
39997,
        0.00040026, 0.          , 0.          ]),
'param_alpha': masked_array(data=[1e-20, 1e-10, 1e-05, 0.1, 1, 10, 100, 20],
        mask=[False, False, False, False, False, False, False, False],
        fill_value='?',
        dtype=object),
'param_max_iter': masked_array(data=[1000000.0, 1000000.0, 1000000.0, 1000000.
0, 1000000.0,
        1000000.0, 1000000.0, 1000000.0],
        mask=[False, False, False, False, False, False, False, False],
        fill_value='?',
        dtype=object),
'params': [{'alpha': 1e-20, 'max_iter': 1000000.0},
{'alpha': 1e-10, 'max_iter': 1000000.0},
{'alpha': 1e-05, 'max_iter': 1000000.0},
{'alpha': 0.1, 'max_iter': 1000000.0},
{'alpha': 1, 'max_iter': 1000000.0},
{'alpha': 10, 'max_iter': 1000000.0},
{'alpha': 100, 'max_iter': 1000000.0},
{'alpha': 20, 'max_iter': 1000000.0}],
'split0_test_score': array([-0.07428915, -0.07428914, -0.07464274, -0.0856067
4, -0.10070146,
        -0.10542019, -0.12704672, -0.10528853]),
'split1_test_score': array([-0.06303485, -0.06303483, -0.06197237, -0.0839421
7, -0.10919212,
        -0.10909901, -0.13063961, -0.10878617]),
'split2_test_score': array([-0.06124805, -0.06124805, -0.0609032 , -0.0706723
6, -0.07527377,
        -0.08466295, -0.12563457, -0.08735557]),
'split3_test_score': array([-0.05335686, -0.05335682, -0.05185672, -0.0667246
, -0.12056159,
        -0.11816192, -0.13079112, -0.11674674]),
'split4_test_score': array([-0.06737639, -0.06737634, -0.06306814, -0.0741482
7, -0.10757998,
        -0.10998106, -0.14368361, -0.1124399 ]),
'mean_test_score': array([-0.06386106, -0.06386104, -0.06248863, -0.07621883,
-0.10266178,
        -0.10546502, -0.13155913, -0.10612338]),
'std_test_score': array([0.00691305, 0.00691305, 0.00726465, 0.00738882, 0.015
1075 ,
        0.01120303, 0.00638593, 0.01012692]),
'rank_test_score': array([3, 2, 1, 4, 5, 6, 8, 7])}
```

```
In [35]: 1 lasso_best_estimator = Lasso(alpha=1e-05,max_iter=1000000.0)
          2
```

```
In [37]: 1 lasso_best_estimator.fit(X_train,y_train)
```

```
Out[37]: Lasso(alpha=1e-05, max_iter=1000000.0)
```

```
In [38]: 1 lasso_best_estimator.score(X_test,y_test)
```

```
Out[38]: 0.7438235254978831
```

```
In [29]: lasso_1e1.score(X_test,y_test)
```

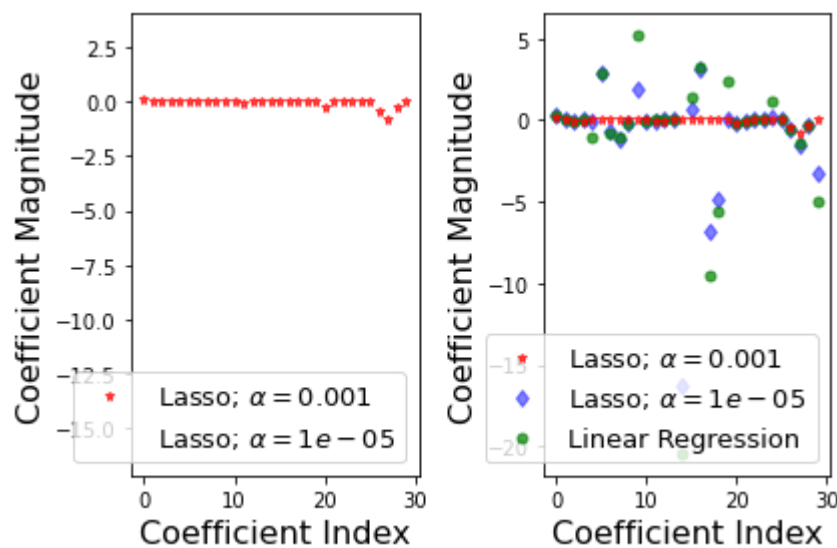
```
Out[29]: 0.7033245684677331
```

**thus after using gridsearchcv we found the best hyperparameters for our Lasso Reg and got a score more than LinearRegression**

```

In [45]: 1 plt.subplot(1,2,1)
2 plt.plot(lasso_1e1.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
3 plt.plot(lasso_best_estimator.coef_,alpha=1e-05,linestyle='none',marker='d',
4
5 plt.xlabel('Coefficient Index',fontsize=16)
6 plt.ylabel('Coefficient Magnitude',fontsize=16)
7 plt.legend(fontsize=13,loc=4)
8
9 plt.subplot(1,2,2)
10 plt.plot(lasso_1e1.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
11 plt.plot(lasso_best_estimator.coef_,alpha=0.5,linestyle='none',marker='d',ma
12
13 plt.plot(lr.coef_,alpha=0.7,linestyle='none',marker='o',markersize=5,color='r')
14 plt.xlabel('Coefficient Index',fontsize=16)
15 plt.ylabel('Coefficient Magnitude',fontsize=16)
16 plt.legend(fontsize=13,loc=4)
17 plt.tight_layout()
18 plt.show()

```



thus we successfully demonstrated how we can improve the accuracy of our model using Normalization techniques like Lasso Regression