

Dhruv Patel

17162121014

Practical 10 - CNN

```
In [2]: ## here we will working with MNSIT dataset
from keras.datasets import mnist
from keras.utils.np_utils import to_categorical
```

```
In [6]: ## Load data
```

```
(X_train,y_train),(X_test,y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-data-sets/mnist.npz>

11493376/11490434 [=====] - 2s 0us/step

```
In [7]: print(X_train.shape) ## 60000 images of 28*28
print(X_test.shape) ## 10000 images of 28*28
```

```
(60000, 28, 28)
```

```
(10000, 28, 28)
```

```
In [8]: print(y_train.shape)
print(y_test.shape)
```

```
(60000,)
```

```
(10000,)
```

```
In [9]: print(y_train[:10])
```

```
[5 0 4 1 9 2 1 3 1 4]
```

```
In [11]: # type(y_train)
```

```
In [14]: X_train = X_train.reshape(60000,28,28,1)
X_test = X_test.reshape(10000,28,28,1)
y_train=to_categorical(y_train)
y_test=to_categorical(y_test)
```

```
In [16]: y_train ## converted into categorical , one hot encoded format
```

```
Out[16]: array([[0., 0., 0., ..., 0., 0., 0.],
               [1., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)
```

```
In [17]: ## train our model
```

```
In [18]: from keras.models import Sequential
         from keras.layers import Flatten,Conv2D,Dense
```

```
In [20]: ## build model
```

```
model = Sequential()
model.add(Conv2D(30,kernel_size=3,activation='relu',input_shape=(28,28,1)))
model.add(Conv2D(15,kernel_size=3,activation='relu'))
model.add(Flatten())
model.add(Dense(10,activation="softmax"))
```

```
In [21]: model.compile(optimizer = 'adam',loss='categorical_crossentropy',metrics=["acc
         uracy"])
```

```
In [22]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 30)	300
conv2d_1 (Conv2D)	(None, 24, 24, 15)	4065
flatten (Flatten)	(None, 8640)	0
dense (Dense)	(None, 10)	86410
Total params: 90,775		
Trainable params: 90,775		
Non-trainable params: 0		

```
In [24]: m1 = model.fit(X_train,y_train,validation_data=(X_test,y_test),epochs=10) ##  
         1st model
```

```
Epoch 1/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0380 - acc  
uracy: 0.9879 - val_loss: 0.0761 - val_accuracy: 0.9811  
Epoch 2/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0289 - acc  
uracy: 0.9912 - val_loss: 0.0933 - val_accuracy: 0.9778  
Epoch 3/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0248 - acc  
uracy: 0.9927 - val_loss: 0.1073 - val_accuracy: 0.9786  
Epoch 4/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0230 - acc  
uracy: 0.9935 - val_loss: 0.1178 - val_accuracy: 0.9786  
Epoch 5/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0193 - acc  
uracy: 0.9941 - val_loss: 0.1572 - val_accuracy: 0.9786  
Epoch 6/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0183 - acc  
uracy: 0.9949 - val_loss: 0.1336 - val_accuracy: 0.9795  
Epoch 7/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0177 - acc  
uracy: 0.9954 - val_loss: 0.1531 - val_accuracy: 0.9805  
Epoch 8/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0145 - acc  
uracy: 0.9962 - val_loss: 0.1695 - val_accuracy: 0.9803  
Epoch 9/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0155 - acc  
uracy: 0.9960 - val_loss: 0.1820 - val_accuracy: 0.9788  
Epoch 10/10  
1875/1875 [=====] - 5s 3ms/step - loss: 0.0126 - acc  
uracy: 0.9969 - val_loss: 0.2013 - val_accuracy: 0.9821
```

```

In [26]: #Plot the train/test accuracy and train/test loss
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np


plt.figure(figsize=(7,3))

plt.subplot(1,2,1)
train_acc = m1.history['accuracy'];
test_acc = m1.history['val_accuracy'];


nepochs = len(train_acc);
sns.lineplot(x=np.arange(1,nepochs+1), y=train_acc, label='Training accuracy'
);
sns.lineplot(x=np.arange(1,nepochs+1), y=test_acc, label='Test accuracy');
plt.xlabel('Epoch');
plt.ylabel('Accuracy');

plt.subplot(1,2,2)
train_loss = m1.history['loss']
test_loss = m1.history['val_loss']
sns.lineplot(x=np.arange(1,nepochs+1), y=train_loss, label='Training loss');
sns.lineplot(x=np.arange(1,nepochs+1), y=test_loss, label='Test loss');
plt.yscale('log')
plt.xlabel('Epochs')
plt.ylabel('Training loss')
plt.tight_layout()

```



```
In [27]: plt.imshow(X_test[17].reshape(28,28), cmap='binary');
```



```
In [32]: # print(model.predict(X_test[[17]]))
print()
print('The number with max probability is')
print(np.argmax(model.predict(X_test[[17]])))
```

The number with max probability is
7

```
In [67]: model2=Sequential()
model2.add(Conv2D(30,kernel_size=(3,3),activation='relu',input_shape=(28,28,1)))
model2.add(Conv2D(64,kernel_size=(3,3),activation="relu"))
model2.add(Conv2D(15,kernel_size=(3,3),activation='relu'))
model2.add(Flatten())
model2.add(Dense(10,activation="softmax"))
```

```
In [68]: model2.compile(optimizer = 'adam',loss='categorical_crossentropy',metrics=["accuracy"])
```

In [69]: `model2.summary()`

Model: "sequential_6"

Layer (type)	Output Shape	Param #
=====		
conv2d_15 (Conv2D)	(None, 26, 26, 30)	300
conv2d_16 (Conv2D)	(None, 24, 24, 64)	17344
conv2d_17 (Conv2D)	(None, 22, 22, 15)	8655
flatten_6 (Flatten)	(None, 7260)	0
dense_6 (Dense)	(None, 10)	72610
=====		
Total params: 98,909		
Trainable params: 98,909		
Non-trainable params: 0		
=====		

In [71]: `m2 = model2.fit(X_train,y_train,validation_data=(X_test,y_test),epochs=10) ##`
2nd model

Epoch 1/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0276 - accuracy: 0.9904 - val_loss: 0.0760 - val_accuracy: 0.9830

Epoch 2/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0228 - accuracy: 0.9928 - val_loss: 0.1091 - val_accuracy: 0.9751

Epoch 3/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0225 - accuracy: 0.9930 - val_loss: 0.1011 - val_accuracy: 0.9795

Epoch 4/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0203 - accuracy: 0.9938 - val_loss: 0.1065 - val_accuracy: 0.9786

Epoch 5/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0163 - accuracy: 0.9953 - val_loss: 0.1131 - val_accuracy: 0.9794

Epoch 6/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0187 - accuracy: 0.9948 - val_loss: 0.1107 - val_accuracy: 0.9802

Epoch 7/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0154 - accuracy: 0.9955 - val_loss: 0.1233 - val_accuracy: 0.9803

Epoch 8/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0153 - accuracy: 0.9956 - val_loss: 0.1524 - val_accuracy: 0.9796

Epoch 9/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0150 - accuracy: 0.9961 - val_loss: 0.1395 - val_accuracy: 0.9797

Epoch 10/10

1875/1875 [=====] - 7s 4ms/step - loss: 0.0121 - accuracy: 0.9968 - val_loss: 0.1610 - val_accuracy: 0.9816

```

In [72]: #Plot the train/test accuracy and train/test loss
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

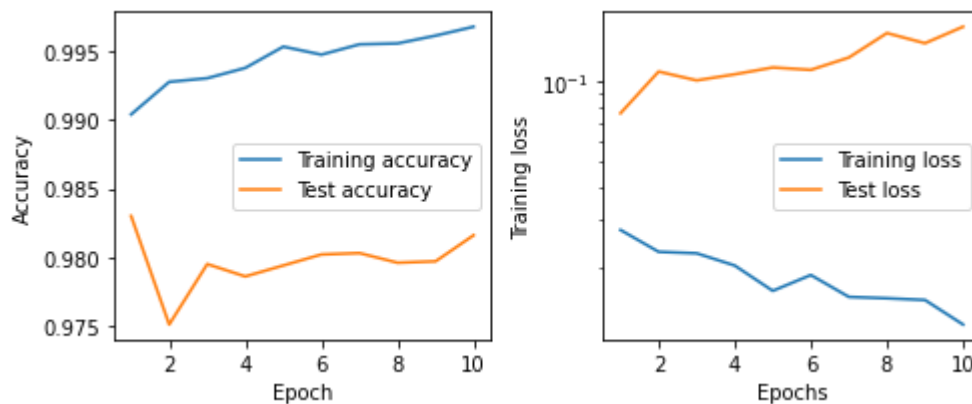
plt.figure(figsize=(7,3))

plt.subplot(1,2,1)
train_acc = m2.history['accuracy'];
test_acc = m2.history['val_accuracy'];

nepochs = len(train_acc);
sns.lineplot(x=np.arange(1,nepochs+1), y=train_acc, label='Training accuracy'
);
sns.lineplot(x=np.arange(1,nepochs+1), y=test_acc, label='Test accuracy');
plt.xlabel('Epoch');
plt.ylabel('Accuracy');

plt.subplot(1,2,2)
train_loss = m2.history['loss']
test_loss = m2.history['val_loss']
sns.lineplot(x=np.arange(1,nepochs+1), y=train_loss, label='Training loss');
sns.lineplot(x=np.arange(1,nepochs+1), y=test_loss, label='Test loss');
plt.yscale('log')
plt.xlabel('Epochs')
plt.ylabel('Training loss')
plt.tight_layout()

```



```

In [ ]: ## thus we tweaked first model and were able to get some slight improvement in
accuracy

```