

```
In [2]: ##import necessary Libraries
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt

from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from sklearn.model_selection import train_test_split
from keras.utils import to_categorical

from keras.models import Sequential
from keras.layers import Dense,Dropout,Embedding,LSTM,GlobalMaxPooling1D,SpatialDropout1D
```

```
In [7]: df_train = pd.read_csv("train.tsv",sep="\t")
df_train.head()
```

```
Out[7]:
```

	PhraseId	SentenceId	Phrase	Sentiment
0	1	1	A series of escapades demonstrating the adage ...	1
1	2	1	A series of escapades demonstrating the adage ...	2
2	3	1	A series	2
3	4	1	A	2
4	5	1	series	2

```
In [8]: df_train.columns
```

```
Out[8]: Index(['PhraseId', 'SentenceId', 'Phrase', 'Sentiment'], dtype='object')
```

```
In [10]: df_train.shape
```

```
Out[10]: (156060, 4)
```

```
In [11]: df_test = pd.read_csv("test.tsv",sep="\t")
df_test.head()
```

```
Out[11]:
```

	PhraseId	SentenceId	Phrase
0	156061	8545	An intermittently pleasing but mostly routine ...
1	156062	8545	An intermittently pleasing but mostly routine ...
2	156063	8545	An
3	156064	8545	intermittently pleasing but mostly routine effort
4	156065	8545	intermittently pleasing but mostly routine

```
In [12]: df_test.shape
```

```
Out[12]: (66292, 3)
```

```
In [13]: df_test.columns
```

```
Out[13]: Index(['PhraseId', 'SentenceId', 'Phrase'], dtype='object')
```


```
In [15]: df_train["Sentiment"].value_counts()## distribution of sentiment
```

```
Out[15]: 2    79582
         3    32927
         1    27273
         4     9206
         0     7072
         Name: Sentiment, dtype: int64
```

```
In [18]: import seaborn as sns
```

```
sns.countplot(df_train["Sentiment"])
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x23f762feb88>
```




```
In [19]: ## cleaning the text
replace_list = {r"i'm": 'i am',
                r"re": ' are',
                r"let's": 'let us',
                r"s": ' is',
                r"ve": ' have',
                r"can't": 'can not',
                r"cannot": 'can not',
                r"shan't": 'shall not',
                r"n't": ' not',
                r"d": ' would',
                r"ll": ' will',
                r"scuse": 'excuse',
                ',': ',',
                '.': '.',
                '!': '!',
                '?': '?',
                '\s+': ' '}
```

```
In [20]: def clean_text(text):  
        text = text.lower()  
  
        for s in replace_list:  
            text = text.replace(s, replace_list[s])  
        text = ' '.join(text.split())  
        return text
```

```
In [21]: X_train = df_train['Phrase'].apply(lambda p: clean_text(p))
```

```
In [23]: phrase_len = X_train.apply(lambda p: len(p.split(' ')))  
max_phrase_len = phrase_len.max()  
print('max phrase len: {}'.format(max_phrase_len))  
plt.figure(figsize = (10, 8))  
plt.hist(phrase_len, alpha = 0.5)  
plt.xlabel('phrase len')  
plt.ylabel('probability')  
plt.grid(alpha = 0.25)
```

max phrase len: 53



```
In [24]: y_train = df_train['Sentiment']
```

```
In [25]: max_words = 8192
tokenizer = Tokenizer(
    num_words = max_words,
    filters = '"#$%&()*+,-./:;<=>@[\\]^_`{|}~'
)
tokenizer.fit_on_texts(X_train)
X_train = tokenizer.texts_to_sequences(X_train)
X_train = pad_sequences(X_train, maxlen = max_phrase_len)

y_train = to_categorical(y_train)
```

```
In [26]: batch_size = 512
epochs = 8
```

```
In [27]: model_lstm = Sequential()
model_lstm.add(Embedding(input_dim = max_words, output_dim=256,input_length =
max_phrase_len))
model_lstm.add(LSTM(256, dropout = 0.3, recurrent_dropout = 0.3))
model_lstm.add(Dense(5, activation = 'softmax'))
model_lstm.compile(
    loss='categorical_crossentropy',
    optimizer='Adam',
    metrics=['accuracy']
)
```

WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

```
In [28]: history = model_lstm.fit(
        X_train,
        y_train,
        validation_split = 0.1,
        epochs = 8,
        batch_size = 512
    )
```

Epoch 1/8

275/275 [=====] - 75s 274ms/step - loss: 1.0325 - accuracy: 0.5882 - val_loss: 0.9996 - val_accuracy: 0.5990

Epoch 2/8

275/275 [=====] - 73s 266ms/step - loss: 0.8266 - accuracy: 0.6682 - val_loss: 1.0029 - val_accuracy: 0.6037

Epoch 3/8

275/275 [=====] - 75s 273ms/step - loss: 0.7793 - accuracy: 0.6844 - val_loss: 1.0165 - val_accuracy: 0.6021

Epoch 4/8

275/275 [=====] - 77s 278ms/step - loss: 0.7473 - accuracy: 0.6960 - val_loss: 1.0330 - val_accuracy: 0.5986

Epoch 5/8

275/275 [=====] - 76s 276ms/step - loss: 0.7202 - accuracy: 0.7059 - val_loss: 1.0587 - val_accuracy: 0.5968

Epoch 6/8

275/275 [=====] - 77s 280ms/step - loss: 0.6954 - accuracy: 0.7153 - val_loss: 1.0653 - val_accuracy: 0.6023


Epoch 7/8

275/275 [=====] - 77s 281ms/step - loss: 0.6734 - accuracy: 0.7231 - val_loss: 1.1138 - val_accuracy: 0.5993


Epoch 8/8

275/275 [=====] - 77s 281ms/step - loss: 0.6550 - accuracy: 0.7298 - val_loss: 1.1239 - val_accuracy: 0.5913

```
In [29]: plt.clf()
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, 'g', label='Training loss')
plt.plot(epochs, val_loss, 'y', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [30]: plt.clf()
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
plt.plot(epochs, acc, 'g', label='Training acc')
plt.plot(epochs, val_acc, 'y', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
In [31]: for key in history.history.keys():
print(key)
```

```
loss
accuracy
val_loss
val_accuracy
```

```
In [32]: X_test = df_test['Phrase'].apply(lambda p: clean_text(p))
X_test.head(5)
tokenizer.fit_on_texts(X_test)
X_test = tokenizer.texts_to_sequences(X_test)
X_test = pad_sequences(X_test, maxlen = max_phrase_len)
```

```
In [33]: y_pred=model_lstm.predict_classes(X_test)
y_pred
```

WARNING:tensorflow:From <ipython-input-33-4a900e69cec3>:1: Sequential.predict_classes (from tensorflow.python.keras.engine.sequential) is deprecated and will be removed after 2021-01-01.

Instructions for updating:

Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

```
Out[33]: array([2, 2, 2, ..., 1, 1, 2], dtype=int64)
```

```
In [34]: batch_size = 512
         epochs = 2
```

```
In [38]: ## model2s with recurrent dropout 0 and recurrent activation ="sigmoid"

model_lstm2 = Sequential()
model_lstm2.add(Embedding(input_dim = max_words, output_dim=256,input_length =
max_phrase_len))
model_lstm2.add(LSTM(256, dropout = 0.3, recurrent_dropout = 0,recurrent_activ
ation="sigmoid"))
model_lstm2.add(Dense(5, activation = 'softmax'))
model_lstm2.compile(
    loss='categorical_crossentropy',
    optimizer='Adam',
    metrics=['accuracy']
)
```

```
In [39]: history2 = model_lstm2.fit(
        X_train,
        y_train,
        validation_split = 0.1,
        epochs = 8,
        batch_size = 512
    )
```

Epoch 1/8

275/275 [=====] - 20s 72ms/step - loss: 1.0334 - acc
uracy: 0.5871 - val_loss: 1.0092 - val_accuracy: 0.5941

Epoch 2/8

275/275 [=====] - 19s 71ms/step - loss: 0.8297 - acc
uracy: 0.6659 - val_loss: 1.0068 - val_accuracy: 0.6002

Epoch 3/8

275/275 [=====] - 21s 76ms/step - loss: 0.7818 - acc
uracy: 0.6846 - val_loss: 1.0119 - val_accuracy: 0.6020

Epoch 4/8

275/275 [=====] - 20s 73ms/step - loss: 0.7484 - acc
uracy: 0.6963 - val_loss: 1.0241 - val_accuracy: 0.6012

Epoch 5/8

275/275 [=====] - 20s 71ms/step - loss: 0.7208 - acc
uracy: 0.7055 - val_loss: 1.0559 - val_accuracy: 0.5953

Epoch 6/8

275/275 [=====] - 20s 72ms/step - loss: 0.6955 - acc
uracy: 0.7154 - val_loss: 1.0681 - val_accuracy: 0.5952

Epoch 7/8

275/275 [=====] - 20s 71ms/step - loss: 0.6740 - acc
uracy: 0.7229 - val_loss: 1.0889 - val_accuracy: 0.5945


Epoch 8/8

275/275 [=====] - 20s 74ms/step - loss: 0.6544 - acc
uracy: 0.7299 - val_loss: 1.1113 - val_accuracy: 0.5981


```
In [41]: plt.clf()
loss = history2.history['loss']
val_loss = history2.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, 'g', label='Training loss')
plt.plot(epochs, val_loss, 'y', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [44]: plt.clf()
acc = history2.history['accuracy']
val_acc = history2.history['val_accuracy']
plt.plot(epochs, acc, 'g', label='Training acc')
plt.plot(epochs, val_acc, 'y', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
In [45]: X_test = df_test['Phrase'].apply(lambda p: clean_text(p))
X_test.head(5)
tokenizer.fit_on_texts(X_test)
X_test = tokenizer.texts_to_sequences(X_test)
X_test = pad_sequences(X_test, maxlen = max_phrase_len)
```

```
In [46]: y_pred2=model_lstm2.predict_classes(X_test)
y_pred2
```

```
Out[46]: array([2, 2, 2, ..., 2, 2, 2], dtype=int64)
```


```
In [48]: df_pred = pd.DataFrame(y_pred2)
```

```
In [50]: df_pred.value_counts()
```

```
Out[50]: 2    42861
3     12122
1       9417
4       1490
0        402
dtype: int64
```

```
In [53]: sns.countplot(df_pred[0])
```

```
Out[53]: <matplotlib.axes._subplots.AxesSubplot at 0x240a50122c8>
```



```
In [ ]:
```